

Diabetes Prediction

December 22, 2022

1 Diabetes Prediction using ML

1.1 Importing important libraries

```
[1]: import numpy as np
import pandas as pd
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC                                #importing support_
    ↪vector regression
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix , classification_report
from sklearn.pipeline import make_pipeline
import pickle
%matplotlib inline
```

2 Data collection and analysis

2.1 Pima dataset

```
[2]: db_df = pd.read_csv('diabetes.csv')
```

```
[3]: #Finding the number of Number of rows and colmns of dataset
db_df.shape
```

```
[3]: (768, 9)
```

```
[4]: #Looking the top 5 data of dataset
db_df.head()
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	

3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[5]: #Getting the statistical measure of data
db_df.describe()
```

```
[5]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count    768.000000    768.000000    768.000000    768.000000    768.000000
mean       3.845052   120.894531     69.105469     20.536458     79.799479
std        3.369578    31.972618     19.355807     15.952218    115.244002
min         0.000000     0.000000     0.000000     0.000000     0.000000
25%         1.000000    99.000000     62.000000     0.000000     0.000000
50%         3.000000   117.000000     72.000000    23.000000    30.500000
75%         6.000000   140.250000     80.000000    32.000000   127.250000
max        17.000000   199.000000    122.000000    99.000000   846.000000
```


	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[6]: #checking for empty value in dataset
db_df.isnull().sum()
```

```
[6]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
[7]: #plotting a heatmap to check for null value
plt.figure(figsize=(15,15))
sn.heatmap(db_df.isnull()) #since we get heatmap of uniform colour so no data is missing
```

[7]: <AxesSubplot:>



```
[8]: #checking data in dataset for true or false data
db_df_shpe = db_df['Outcome'].value_counts()
print('The total data not having diabetes:--{}\n\The total data having diabetes:
↳--{}'.format(db_df_shpe[0],db_df_shpe[1]))
```

The total data not having diabetes:-500
The total data having diabetes:-268

```
[9]: #Getting more analysis of our data
db_df.groupby('Outcome').mean()
```

```
[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
Outcome					
0	3.298000	109.980000	68.184000	19.664000	68.792000
1	4.865672	141.257463	70.824627	22.164179	100.335821

	BMI	DiabetesPedigreeFunction	Age
Outcome			
0	30.304200	0.429734	31.190000
1	35.142537	0.550500	37.067164

```
[10]: #seprating the data and labels
X = db_df.drop(['Outcome'],axis=1)
y = db_df['Outcome']
```

```
[11]: X
```

```
[11]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4


	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
763	0.171	63

764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[768 rows x 8 columns]

3 Creating a correlational matrix

```
[12]: plt.figure(figsize=(15,15))
heatmap = sn.heatmap(db_df.corr(),annot=True,cmap="YlGnBu");
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':12}, pad=12);
```



3.1 Data Standardization

```
[13]: scaler =StandardScaler()
      X_standardised = scaler.fit_transform(X)
```

```
[14]: X_standardised
```

```
[14]: array([[ 0.63994726,  0.84832379,  0.14964075, ...,  0.20401277,
               0.46849198,  1.4259954 ],
             [-0.84488505, -1.12339636, -0.16054575, ..., -0.68442195,
               -0.36506078, -0.19067191],
             [ 1.23388019,  1.94372388, -0.26394125, ..., -1.10325546,
               0.60439732, -0.10558415],
             ...,
             [ 0.3429808 ,  0.00330087,  0.14964075, ..., -0.73518964,
               -0.68519336, -0.27575966],
             [-0.84488505,  0.1597866 , -0.47073225, ..., -0.24020459,
               -0.37110101,  1.17073215],
             [-0.84488505, -0.8730192 ,  0.04624525, ..., -0.20212881,
               -0.47378505, -0.87137393]])
```

3.2 Train test split

```
[15]: X_train ,X_test,y_train,y_test = train_test_split(X_standardised,y,test_size=0.
        ↪2,random_state=12,stratify=y)
```

```
[16]: print('The original data shape is {}. Test data shape {} and train data shape_
        ↪is {}'.format(X.shape,X_train.shape,X_test.shape))
```

The original data shape is (768, 8). Test data shape (614, 8) and train data shape is (154, 8)

3.3 Selecting the best possible model with hyperparameters between logistic_regression and support vector classifier.

```
[17]: model_params = {
        'logistic_regression' :{
            'model' : LogisticRegression(),
            'params' :{
                'penalty':['l1', 'l2', 'elasticnet', None],
                'C':[-7,1e-2,0,1,2,3,4,5,6,7,8,9,10,20,30,40,50],
                'max_iter':[10,50,100,200,300,500],
                'tol':[1e-5,1e-4,1e-6,1e-8]
            }
        },
        'SVC':{
            'model' : SVC(),
            'params':{
                'gamma':['auto','scale'],
                'C':[-7,1e-2,0,1,2,3,4,5,6,7,8,9,10,20,30,40,50],
                'kernel':['linear', 'poly', 'rbf', 'sigmoid'],
                'coef0':[0.0,0.5,0.7,0.9,1.0,2.0]
            }
        }
    }
```

```
}
```

```
[18]: scores = []
      for model_name, mp in model_params.items():
          clf = GridSearchCV(mp['model'], mp['params'], cv = 5, return_train_score=False)
          clf.fit(X_train, y_train)
          scores.append({
              'model': model_name,
              'best_score': clf.best_score_,
              'best_params': clf.best_params_,
          })
```

```
C:\Users\01abn\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\01abn\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\01abn\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\Users\01abn\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
```



```

0.7100893 0.67093163 0.76699987 0.70194589 0.70846328 0.64647474
0.76699987 0.70194589 0.7100893 0.6497401 0.76699987 0.69706784
0.70846328 0.65621751 0.76699987 0.69706784 0.7100893 0.63833133
0.76699987 0.70031987 0.70846328 0.66613355 0.76699987 0.69869386
0.7100893 0.66775956 0.76699987 0.69544182 0.70846328 0.66775956
0.76699987 0.69706784 0.7100893 0.66450753 0.76699987 0.69706784
0.70846328 0.65631081 0.76699987 0.69869386 0.7100893 0.64167666
0.76699987 0.69540184 0.71985872 0.67255764 0.76699987 0.6937625
0.71660669 0.6709183 0.76699987 0.69541517 0.71985872 0.64484873
0.76699987 0.69541517 0.71660669 0.64484873 0.76699987 0.69540184
0.71985872 0.65133946 0.76699987 0.69540184 0.71660669 0.67098494
0.76699987 0.69702786 0.71985872 0.66775956 0.76699987 0.69702786
0.71660669 0.66448087 0.76699987 0.70027989 0.71985872 0.65312542
0.76699987 0.69702786 0.71660669 0.65475143 0.76699987 0.69868053
0.71985872 0.63681194 0.76699987 0.6986672 0.71660669 0.63843796]
warnings.warn(

```

```
[19]: t_df = pd.DataFrame(scores,columns=['model','best_score','best_params'])
```

```
[20]: print(t_df)
```

```

              model  best_score \
0  logistic_regression    0.767013
1                SVC      0.768626

              best_params
0  {'C': 1, 'max_iter': 10, 'penalty': 'l2', 'tol...
1  {'C': 1, 'coef0': 0.0, 'gamma': 'auto', 'kerne...

```

```
[31]: clf.best_score_
```

```
[31]: 0.7686258829801413
```

4 Therefore on comparing the performance of both the model we conclude that svc work better

4.0.1 Model evaluation

Accuracy score

```
[21]: y_test_predict = clf.predict(X_test)
y_train_predict = clf.predict(X_train)
res1= accuracy_score(y_test_predict,y_test)
res2 = accuracy_score(y_train_predict,y_train)
```

```
[22]: print('Accuracy score on test is {} and train data for our model is {}'.
        format(res1,res2))
```

```


Accuracy score on test is 0.7857142857142857 and train data for our model is
0.7687296416938111

```

```
[23]: #creating the confusion matrix for train data checking the accuracy of our model
cm= confusion_matrix(y_true = y_train,y_pred=y_train_predict)

plt.figure(figsize=(10,7))
sn.heatmap(cm,annot=True,fmt='d')
plt.xlabel('Predected')
plt.ylabel('Truth')
```

```
[23]: Text(69.0, 0.5, 'Truth')
```



4.1 Making a predictive system

```
[29]: input_data = (5,187,76,27,207,43.6,1.034,53)

# changing the input_data to numpy array
input_data = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data = input_data.reshape(1,-1)
```

```

#standarised the data
input_data_std = scaler.transform(input_data)
print(input_data_std)

prediction = clf.predict(input_data_std)
if prediction[0] == 1:
    print('The patient has Diabetes')
else:
    print('The patient has not Diabatese')

```

```

[[0.3429808  2.06891246 0.35643175 0.40544544 1.10446888 1.47320522
  1.69768028 1.68125866]]

```

The patient has Diabetes

C:\Users\01abn\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names

```
warnings.warn(
```

4.2 Saving the model

```

[25]: #Saving the scaler and model
filename = 'diabetes_model.sav'
pickle.dump(clf, open(filename, 'wb'))
scalerfile = 'scaler.sav'
pickle.dump(scaler, open(scalerfile, 'wb'))

```

```

[26]: # loading the saved model
load_model = pickle.load(open('diabetes_model.sav', 'rb'))
load_scaler = pickle.load(open('scaler.sav', 'rb'))

```

```

[40]: input_data = (5,104,74,0,0,28.8,0.153,48)

# changing the input_data to numpy array
input_data = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data = input_data.reshape(1,-1)

#standarised the data
print('before input data',input_data)
input_data_std = load_scaler.transform(input_data)
print('after std input data',input_data_std)
print()
prediction = clf.predict(input_data_std)
if prediction[0] == 1:

```

```
print('The patient has Diabetes')
else:
    print('The patient has not Diabatese')
```

```
before input data [[ 5.    104.    74.     0.     0.    28.8    0.153  48.
]]
```

```
after std input data [[ 0.3429808 -0.5287506  0.25303625 -1.28821221
-0.69289057 -0.40519961
-0.96304428  1.2558199 ]]
```

The patient has not Diabatese

C:\Users\01abn\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names

```
warnings.warn(
```

```
[38]: for col in db_df:
        print(col)
```

Pregnancies

Glucose

BloodPressure

SkinThickness

Insulin

BMI

DiabetesPedigreeFunction

Age

Outcome

```
[ ]:
```

```
[ ]:
```